

A few Snowflake best practices:

- **Choose a sensible warehouse size** – consider the task you are performing.
- **Have separate warehouses** for different tasks – data loading, analytics, etc.
- **Enable auto-suspend** on warehouses, so when queries are done, they turn off.
- **Enable auto-resume** on warehouses, so they start back up, when needed. (Similar to above.)
- **Apply resources monitors to warehouses** – in order to control costs and avoid spikes.
- **Partition staged data** into multiple, compressed files (100-250MB) and load in parallel. *
- **Only select the columns you need** – this is important with columnar data stores. ^

* *The number of load operations that run in parallel cannot exceed the number of data files to be loaded."*

^ *You can avoid columns that are not relevant.*